



UNIVERSITY OF
SAN FRANCISCO

CHANGE THE WORLD FROM HERE

Support Vectors

Machine Learning



Motivation

- Can you recognise handwritten digits?




Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

- Let's first discuss a popular classifier: Support Vector Machines






Why Support Vectors?

- Pros
 - Memory and computationally inexpensive because it uses a subset of data
 - Easy to interpret results
 - Makes few errors
- Cons
 - Sensitive to configuration and outliers — needs small-ish dataset
 - Natively sensitive to class imbalances
 - Inherently expects pairs of classes
- Other stuff
 - Popular implementations in scikit-learn and (highly optimized) libsvm
 - Some division of opinion: considered old by some, powerful by others



Drawing a Line

- Assume: linearly separable classes
- Goal: Find a *hyperplane* (decision boundary with $p-1$ dimensions)
 - Of the infinite possible hyperplanes, find the one with the largest width
 - Maximize the margin: gap between hyperplane and nearest observation





Class Labels & Hyperplane

- Two classes (1, -1) — assumes points equidistant from margin
- Hyperplane defined by:

$$f(x) = \omega^T x + b$$

... where:

ω^T is a vector of weights


x is the input data

b is an offset weight

- Goals:
 - Find support vectors — “extreme” instances close to frontier but not in (or past) margin
 - $f(x) \geq 1$ for all support vectors in class 1; $f(x) \leq -1$ for all support vectors in class 2
 - Maximise the margin $[z = f(x) / \|\omega^T\|]$ by minimising weight vector (ω^T)




Two Points \Rightarrow Hyperplane (1)



- Goal: find the values for ω , b
- Using 2 points:
 - (1, 1) to class 1 $\Rightarrow f(1, 1) = -1$
 - (2, 3) to class 2 $\Rightarrow f(2, 3) = 1$
- Plane is equidistant from points:
 - $\omega = (2, 3) - (1, 1) = (a, 2a)$
 - $f(1, 1) = -1 = a + 2a + b$
 - $f(2, 3) = 1 = 2a + 6a + b$
- So b has the value $[1 - 8a]$



Two Points \Rightarrow Hyperplane (2)



- Knowing:
 - b has the value $[1 - 8a]$
 - $f(1, 1) = a + 2a + b = -1$
- We can determine ω :
 - Point $(1, 1) = a + 2a + 1 - 8a = -1$
 - ... therefore $a = \frac{2}{5}$
 - $\omega = (\frac{2}{5}, \frac{4}{5})$
- We can determine b :
 - $a + 2a + b = -1$
 - $\frac{2}{5} + \frac{4}{5} + b = -1$
 - ... so $b = -11/5$




More Than Two Points?

- Use support vectors (α = “alphas”)
 - ... i.e. instances closest to the margin
 - ... to maximise margin, given by:

$$\arg \max_{w,b} \left\{ \min_n (label \cdot (w^T x + b)) \cdot \frac{1}{\|w\|} \right\}$$

- Problem: this is difficult to optimise
- Solution:
 - Constrained Optimisation
 - Set $[label * w^T x + b]$ to 1^*

$$\max_{\alpha} \left[\sum_{i=1}^m \alpha - \frac{1}{2} \sum_{i,j=1}^m label^{(i)} \cdot label^{(j)} \cdot a_i \cdot a_j \langle x^{(i)}, x^{(j)} \rangle \right]$$





Slack Variables

- Problems:
 - What if an instance appears within the margin?
 - Worse, what if an instance appears on the wrong side of the margin? (Data is not always linearly separable)
- Solution: Further constrain the solution to exclude these problem points

$$c \geq \alpha \geq 0, \text{ and } \sum_{i=1}^m \alpha_i \cdot \text{label}^{(i)} = 0$$

- “c” controls weighting between two goals:
 - Ensuring most instances have a margin of 1.0 or greater
 - Making the margin large




SMO

- Platt, 1996: SMO (Sequential Minimal Optimization)
 - Was it conceived before?
 - Breaks optimisation into several smaller problems, each with relatively easy solutions
- Basic algorithm
 - Choose a (the best?) pair of alphas
 - Alphas must be outside margin
 - Alphas cannot be “clamped” or “bounded”
 - Determine b
 - Use alphas & b to find weights
 - Harrington’s implementation is on github, [here](#)
- Advanced
 - Heuristics to select good candidates for alpha pairs
 - Classes for holding data structures



Lab 2


- Get Harrington's SMO implementation (**Non-Kernel version**)
 - <https://github.com/pbharrin/machinelearninginaction3x/blob/master/Ch06/svmMLiA.py>
 - Note that Harrington conflates loading the data (loadDataSet) with the classifier; you should NOT do this
- Get the data
 - Download from [linearly_separable.csv](#)
 - Create a figure to ensure classes are linearly separable, eg:
- Change Harrington's code to a classifier
 - Implement "fit" and "predict"
 - Use defaults for C: 1.0; toler: 0.001; maxIter: 50
 - Ignore





Nonlinearities and Kernels

- Problem: linear separation
 - Sometimes class separations would give poor performance
 - Tuning “C” would not help in these cases
- Solution:
 - Move to a higher-dimension space
 - Scalar from two (feature) vectors $\langle x_i, x_j \rangle$
 - Linear support vector classifier →
 - Adds non-linear kernels (Polynomial and RBF = radial basis function)
 - Example: d-dimensional polynomials





SVM vs. Logistic Regression

- Options for $K > 2$ classes
 - OVA (one vs. all): fit K different 2-class SVM classifiers; classify x^* to the class which generates the max score
 - OVO: (one vs. one): Fit all K choose 2 pairwise classifiers; classify x^* to the class which wins the most pairwise competitions
- When classes are (nearly) separable, SVM generally performs better
- When classes are not linearly separable, both behave similarly
 - LR must have *ridge penalty*
 - Need to estimate probabilities? LR is the better choice
- SVMs are popular for non-linear boundaries
 - Implemented as kernels
 - Can also apply kernels to LR or LDA... but don't



Next Time

- Handwritten digits practical