



UNIVERSITY OF
SAN FRANCISCO

CHANGE THE WORLD FROM HERE

Random Forest (and Boosting)

Machine Learning



Recall Decision Trees

- Use a metric to determine how to divide the data into regions
 - Metrics: entropy, gini (eg. impurity)
 - Recursive (Binary) Splitting: continue dividing data into regions until “convergence”
- Advantages
 - Explanatory power
 - Elegant handling of irrelevant features
- Disadvantages
 - Overfitting, perhaps due to high variance
 - Mediocre performance, perhaps due to greedy nature of algorithm



Random Forest Overview

- Build multiple `decision_tree` instances
 - Number of trees is one hyperparameter
 - More trees = more accurate hypotheses
 - Typical to see: 10, 30, 100
- For fit:
 - Build multiple independent decision trees
 - Randomly select data for each tree
 - Randomly select features for each tree
- For predict:
 - Get independent predictions from each tree
 - Take the majority (classification) or mean (regression) of all trees



Advantages & Disadvantages

- Advantages
 - Handles missing data well
 - Handles large amounts of data very well (i.e. either large n , large p or both)
 - Handles small data well
 - Does not overfit*
- Disadvantages
 - Cannot predict outside training targets for regression
 - May overfit on noisy datasets, especially if small
 - Drives statisticians (and other theory-based people) insane
- Other random stuff
 - General purpose algorithm: can generate prediction contours of arbitrary shape
 - Used for body part prediction in MicroSoft Xbox Kinect?




Handling High Variance

- Central Limit Theorem
 - Distribution of repeated sampling of (training) data is a normal distribution
 - True even if data has high variance
- Multiple hypotheses are better than one
 - Increasing sample size helps
 - Increasing number of samples helps
 - Applies to regression as well as classification
- Bagging: Bootstrap Aggregation
 - Sample (with replacement) to create B different bootstrapped datasets

$$\hat{f}_{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

- ... where $\hat{f}^{*b}(x)$ = the prediction point at x for the Bth training set



By Source, Fair use,
<https://en.wikipedia.org/w/index.php?curid=25461520>



Bagging on Animals

animals

| name | convering | eggs | lives_in | oxygenates | legs | tail | category |
|----------|-------------|------|----------|------------|------|------|----------|
| aardvark | hair | 0 | ground | air | 4 | 0 | mammal |
| antelope | hair | 0 | ground | air | 4 | 1 | mammal |
| bass | scales | 1 | water | water | 0 | 1 | fish |
| carp | scales | 1 | water | water | 0 | 1 | fish |
| chicken | feathers | 1 | ground | air | 2 | 1 | bird |
| scorpion | exoskeleton | 0 | ground | air | 8 | 1 | invert |
| starfish | skin | 1 | water | water | 5 | 0 | invert |
| swan | feathers | 1 | air | air | 2 | 1 | bird |

- Example:
 - Randomly generated “n” samples
 - Set 1:

[4, 0, 3, 3, 6, 6, 7, 1]
 - Not shown:
sets 2



Handling Greedy Algorithm

- During “fit”, remove one or more features from consideration
- Randomly decide on features to be removed
 - The scikit-learn implementation considers \sqrt{p} features

From
<http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>

max_features : int, float, string or None, optional (default="auto")

The number of features to consider when looking for the best split:

- If int, then consider max_features features at each split.
- If float, then max_features is a percentage and $\text{int}(\text{max_features} * n_features)$ features are considered at each split.
- If "auto", then $\text{max_features} = \sqrt{n_features}$.
- If "sqrt", then $\text{max_features} = \sqrt{n_features}$ (same as "auto").
- If "log2", then $\text{max_features} = \log_2(n_features)$.
- If None, then $\text{max_features} = n_features$.

Note: the search for a split does not stop until at least one valid partition of the node samples is found, even if it requires to effectively inspect more than `max_features` features.



Random Feature Selection

animals

| name | convering | eggs | lives_in | oxygenates | legs | tail | category |
|----------|-------------|------|----------|------------|------|------|----------|
| aardvark | hair | 0 | ground | air | 4 | 0 | mammal |
| antelope | hair | 0 | ground | air | 4 | 1 | mammal |
| bass | scales | 1 | water | water | 0 | 1 | fish |
| carp | scales | 1 | water | water | 0 | 1 | fish |
| chicken | feathers | 1 | ground | air | 2 | 1 | bird |
| scorpion | exoskeleton | 0 | ground | air | 8 | 1 | invert |
| starfish | skin | 1 | water | water | 5 | 0 | invert |
| swan | feathers | 1 | air | air | 2 | 1 | bird |

- Example:
 - Randomly generated "sqrt(p)" features
 - Estimator 1:
[3, 0]



Random Forest — fit()

```
def fit(X, Y, num_trees)

    tree_list = create_list(num_trees)  # decision_tree list

    for t in tree_list:

        subsample_x, subsample_y = subsample(X, Y) # Bagging

        feature_list = sample_of_features(X) # Random features

        t.fit(subsample_x, subsample_y, feature_list)
```



Random Forest — predict()

```
def predict(X):  
  
    from collections import defaultdict  
  
    hypothesis_list = [t.predict(X) for t in tree_list]  
  
    counts = defaultdict(int)  
  
    for h in hypothesis_list:  
        counts[h] += 1  
  
    return sorted(counts.items(), reverse=True, key=lambda  
tup: tup[1])[:len(tree_list)][0][0]
```



Odds & Ends

- Classically, the random forest algorithm does prune trees
 - Practically, this can be driven by the data
 - In sklearn, this is a hyperparameter (max_depth)
- In sklearn, we can determine feature importances on fitted data, eg:

```
clf.fit(X, Y)
```

```
importances = clf.feature_importances_
```

- The BaggingClassifier class in sklearn is a closely-related class
- Ensemble learning
 - One of many techniques
 - Several weak learners (performing slightly better than chance) may create a strong learner



Boosting Overview

- Observation
 - What if we train a learner (classifier) on some data then have it predict the training data?
 - Answer: The classifier makes errors even on the training set
- Can you learn from mistakes?
 - Create several (sequential) learners
 - Each learner focuses on the previous learner's errors



AdaBoost

- Adaptive Boosting
- Algorithm overview
 - Create “B” data subsets using bagging and “B” learners, one for each bag
 - 1: Train a learner on its bag
 - 2: Test the learner against its own (training set) bag
 - 3: Add probability to training instances which were misclassified so that it's more likely to be selected for subsequent bags
 - 4: Repeat (1) for next bag



Exercise — Occupancy Detection

- Download the Occupancy Detection dataset
 - <https://archive.ics.uci.edu/ml/datasets/Occupancy+Detection+>
 - Target (last column) = 1 (occupied) vs. 0 (unoccupied)
 - Features = Everything except target
 - Ignore instance number and date columns (for now)
- Without finding optimal hyperparameters, try three sklearn classifiers
 - DecisionTreeClassifier (sklearn.tree)
 - RandomForestClassifier (sklearn.ensemble)
 - AdaBoostClassifier (sklearn.ensemble)
- Answer the following:
 - What is the performance vs. datatest?
 - What is the performance vs. datatest2?
 - Can you improve the performance for each?



Next Time

- Hands-on with Random Forest (and AdaBoost / XGBoost?)
- Perceptron & Multi-Layer Perceptron