



UNIVERSITY OF
SAN FRANCISCO

CHANGE THE WORLD FROM HERE


Neural Networks

Machine Learning



Historically...

- Early attempts to model brain (perceptron)
 - Inputs (x) to a function: 0, 1
 - Output (also 0, 1) depends on $f(x)$ based on:
 - Internal weights
 - Internal threshold for activation
- Other attempts: understand Artificial Intelligence possibilities and limits
- Metaphors:
 - Deciding on what to eat at Bon Appetit (considering budget, preferences, etc.)
 - Computer circuits
- Many names for similar items
 - Basic: (multi-layer) perceptron, (artificial) neural network
 - Closely related: dendritic computation?






Perceptrons to Neurons

- Input is multiple vectors of Boolean propositions
- Output = 0 or 1:
 - $\sum_j w_j x_j$
 - Sum passed through an *activation function*, f , (eg. sigmoid)

$$f(x) = \frac{1}{1 + e^{-x}}$$



- Vector product assumed and threshold rewritten as bias (b), so:

$$\text{output} = f [(w \cdot x) + b]$$

- Training a perceptron involves setting weights
- Perceptron to neuron: input and output values do not have to be Boolean



Activation Functions

- Sigmoid may be the most common activation function

- Others:


- Tanh

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

- ReLU ("Rectifier" / "soft plus")

$$f(x) = \log(1 + \exp x)$$


- Others (maxout, leaky, etc.)





The XOR problem

- Early problem in neural networks
- Consider the XOR circuit:
 - Two inputs: x_0, x_1
 - Output:
 - 1 if EITHER x_0 or x_1 has the value 1
 - 0 if x_0 and x_1 have the same value (either 1 or 0)
- Need multiple decisions to model this




<http://toritris.weebly.com/perceptron-5-xor-how--why-neurons-work-together.html>



Designing Networks

- A typical neural network contains
 - Input layer (exactly 1)
 - Hidden layer(s): 0, 1 or more layers
 - Output layer (exactly 1)
- Layer contents
 - Input: As many neurons as features
 - Output: As many neurons as classes*
 - Hidden: No limits
- Feed forward network
 - Output from one layer becomes input to next
 - Network is *fully connected*: all outputs from one layer are inputs to all neurons in next
 - Weights for each neuron provide differentiation among neurons



<https://hackernoon.com/challenges-in-deep-learning-57bbf6e73bb>



Advantages & Disadvantages

- Advantages
 - In theory: some network can learn any function
 - No need to derive features... just feed the network the “raw” signal
 - Unbelievable performance
- Disadvantages
 - No one REALLY understands how they work, so drive statisticians crazy
 - Require a lot of computational resources (GPU?) to train
 - Can overfit



A Neural Network Implementation

```
class Network(Classifier):  
  
    def __init__(self, sizes, epochs, batch_size, learning_rate):  
        self.num_layers = len(sizes)  
        self.sizes = sizes          # Not shown: setting other class data  
  
        # Randomly set biases for all non-input neurons  
        self.biases = [np.random.randn(y, 1) for y in sizes[1:]]  
  
        # Randomly set weights for each input/output pair  
        self.weights = [np.random.randn(y, x) for x, y in  
                        zip(sizes[:-1], sizes[1:])]
```

- Example use:

```
dnn = Network([2, 2, 1])
```




Training: SGD, Backpropagation

- [Stochastic] gradient descent — recall:
 - Calculate *cost function* = some difference between targets and hypotheses
 - Cost function is often SSE
 - Gradient descent minimises cost function by (slowly) moving the decision boundary in the direction of the falling gradient
 - Stochastic: make changes with part of the training data (“batch”)
- Backpropagation (concept)
 - Propagate errors to all layers of the network for each batch
 - Each neuron updates its decision boundary



Backpropagation: Partial Code

```
def backprop(self, x, y):  
  
    # Not shown = declaration of layer-by-layer errors  
    # Not shown = call to get individual neuron activations & hypotheses  
    delta = self.cost_detivative(activations[-1], y) *  
    sigmoid_prime(hyp[-1])  
    bias_error[-1] = delta  
    weight_error[-1] = np.dot(delta, activations[-2].transpose())  
  
    # Step backward in the network and spread errors to all layers (not  
    input)  
    for layer in xrange(2, self.num_layers):  
        z = zs[-layer]  
        sigmoid_prime = sigmoid_prime(z)  
        bias_error[-layer] = delta  
        weight_error[-layer] =  
np.dot(delta, activations[-layer-1].transpose())  
    return (bias_error, weight_error)
```



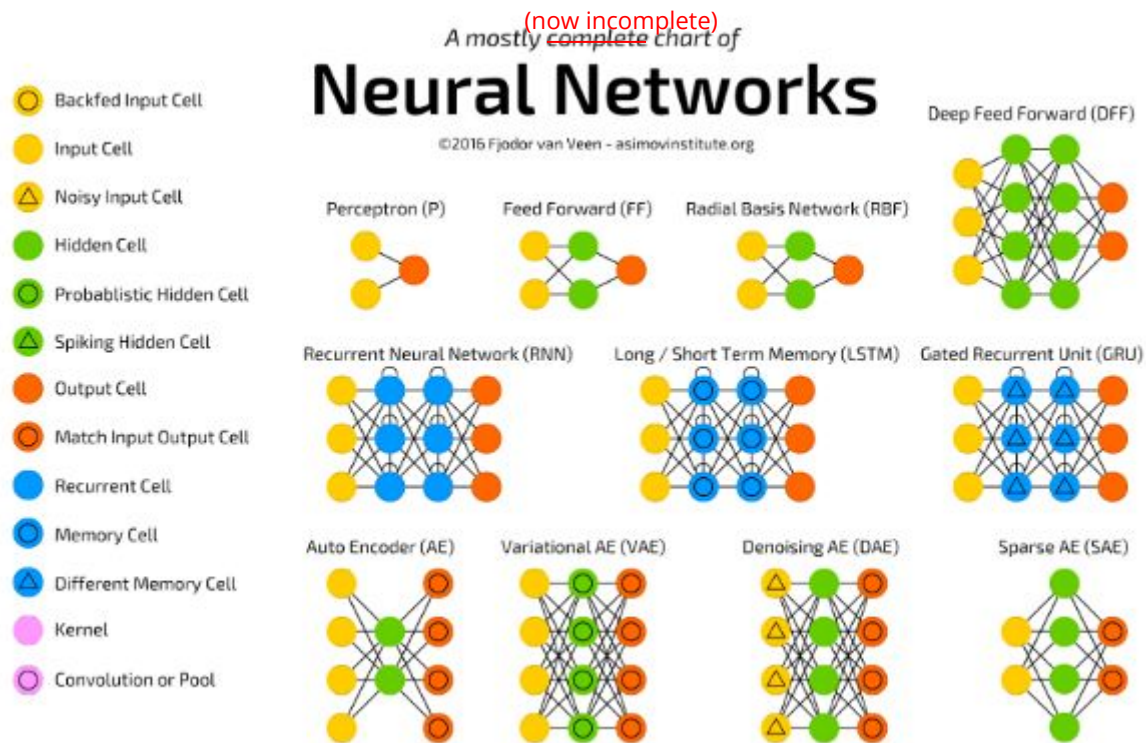
Hands-on (Practical)

- Download MNIST dataset from Kaggle
 - <https://www.kaggle.com/c/digit-recognizer/data>
 - Only need “train.csv” since “test” does not have labels
- Recall model performance from earlier SVM experiments
 - Test set is final 100 items
 - Expected SVM accuracy: $0.78 \leq x \leq 0.91$ (changes with random seed)
- Design a neural network for this
 - Try the following:

```
from sklearn.neural_network import MLPClassifier  
mlp = MLPClassifier(hidden_layer_sizes=(100, 10))
```
 - Which items did it get wrong? (Are they also hard for you to determine?)
 - Experiment with parameters to raise performance



Going Deeper





Tools for Larger Networks



Keras

PYTORCH



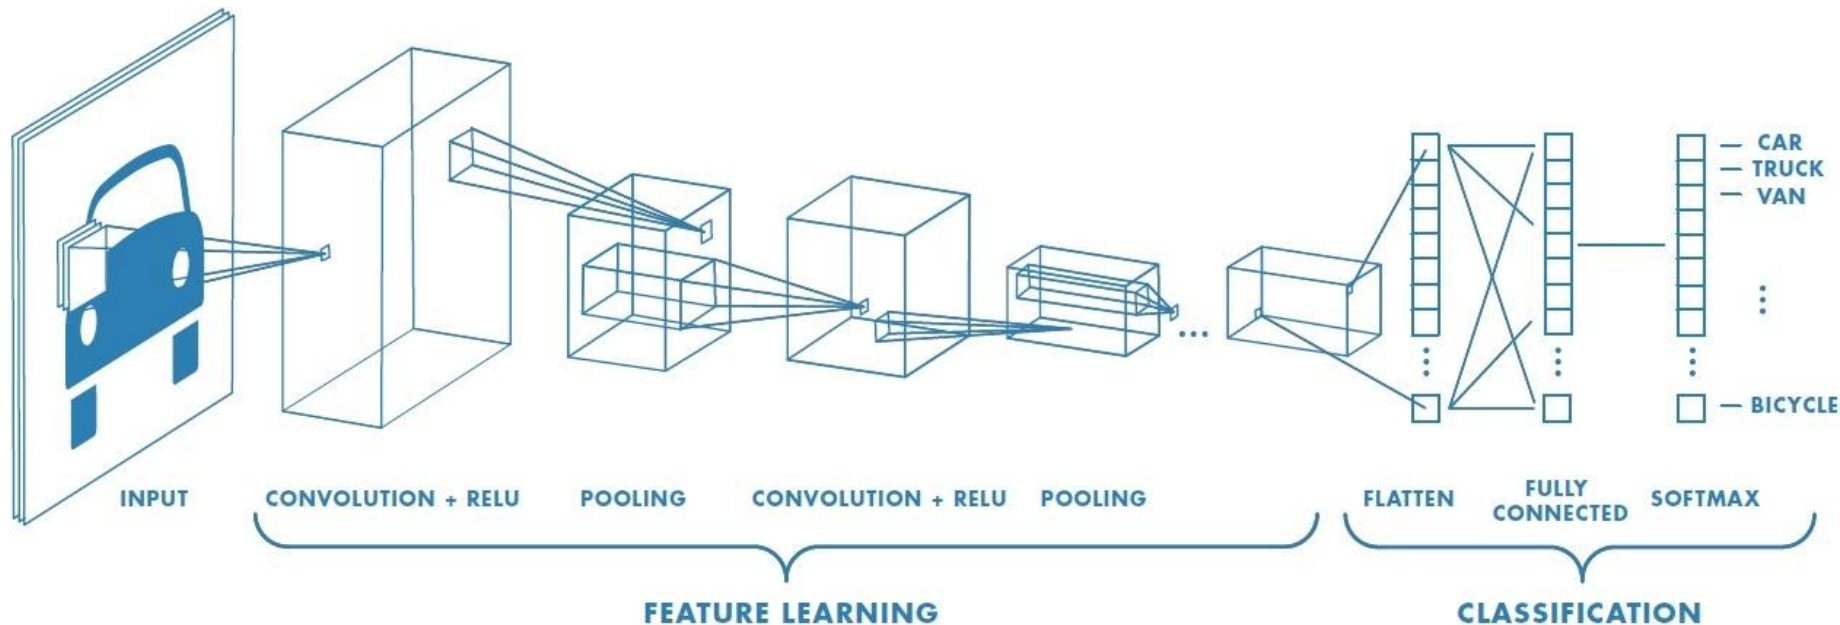
theano



TensorFlow



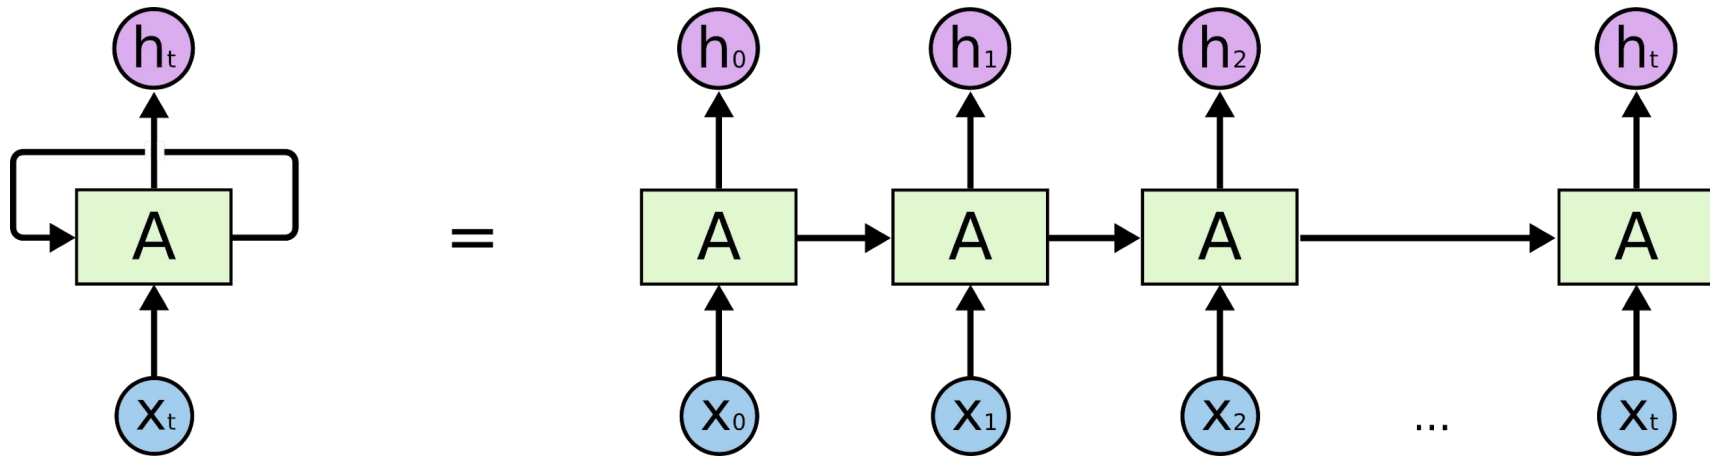
Convolutional Neural Networks





Long Short Term Memory (LSTM)

- Output of a layer becomes (additional) input



- Good for language, speech, video

<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>



Odds and Ends

- Dropout
 - Neurons in the network randomly select some input weights to be 0
 - Forces network to learn more robust features
 - Reduces overfitting — similar to Random Forest, Boosting, etc.
- GPUs are standard equipment for neural networks
 - For each epoch, each connection between two neurons is a matrix calculation
 - CPUs are not optimised to run many tiny calculations
- Data determines outcomes
 - A neural network (any ML algorithm) will only be as good as the data it's trained on
 - Collect data with care!



Next Time

- Thursday:
 - Classifying smile types (Leon Wang)
 - Classification Summary & Review
- Next week: Clustering (Paul Intrevado)
- Midterm II